

### Addition of Vectors using CUDA Code:

```
%%cu
#include "cuda_runtime.h"
#include "device_launch_parameters.h"
#include <stdio.h>

__global__ void addKernel(int* c, const int* a, const int* b, int size) {
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if (i < size) {
        c[i] = a[i] + b[i];
    }
}

// Helper function for using CUDA to add vectors in parallel.
void addWithCuda(int* c, const int* a, const int* b, int size) {
    int* dev_a = nullptr;
    int* dev_b = nullptr;
    int* dev_c = nullptr;

    // Allocate GPU buffers for three vectors (two input, one output)
    cudaMalloc((void**)&dev_c, size * sizeof(int));
    cudaMalloc((void**)&dev_a, size * sizeof(int));
    cudaMalloc((void**)&dev_b, size * sizeof(int));

    // Copy input vectors from host memory to GPU buffers.
    cudaMemcpy(dev_a, a, size * sizeof(int), cudaMemcpyHostToDevice);
    cudaMemcpy(dev_b, b, size * sizeof(int), cudaMemcpyHostToDevice);

    // Launch a kernel on the GPU with one thread for each element.
    // 2 is number of computational blocks and (size + 1) / 2 is a number of threads in a block
    addKernel<<<2, (size + 1) / 2>>>(dev_c, dev_a, dev_b, size);

    // cudaDeviceSynchronize waits for the kernel to finish, and returns any errors encountered during
    //the launch.
    cudaDeviceSynchronize();

    // Copy output vector from GPU buffer to host memory.
    cudaMemcpy(c, dev_c, size * sizeof(int), cudaMemcpyDeviceToHost);

    cudaFree(dev_c);
    cudaFree(dev_a);
    cudaFree(dev_b);
}

int main(int argc, char** argv) {
    const int arraySize = 5;
    const int a[arraySize] = { 1, 2, 3, 4, 5 };
```

```

const int b[arraySize] = { 10, 20, 30, 40, 50 };
int c[arraySize] = { 0 };
addWithCuda(c, a, b, arraySize);
printf("{1, 2, 3, 4, 5} + {10, 20, 30, 40, 50} = {%d, %d, %d, %d, %d}\n", c[0], c[1], c[2], c[3], c[4]);
cudaDeviceReset();
return 0;
}

```

### **Addition of Vectors using CUDA Output:**

$$\{1, 2, 3, 4, 5\} + \{10, 20, 30, 40, 50\} = \{11, 22, 33, 44, 55\}$$

### **Matrix Multiplication using CUDA C Code:**

```

%%cu
#include<stdio.h>
#include<cuda.h>
#define row1 2 /* Number of rows of first matrix */
#define col1 3 /* Number of columns of first matrix */
#define row2 3 /* Number of rows of second matrix */
#define col2 2 /* Number of columns of second matrix */
__global__ void matproduct(int *l, int *m, int *n){
    int x = blockIdx.x;
    int y = blockIdx.y;
    int k;
    n[col2 * y + x] = 0;
    for(k = 0; k < col1; k++){
        n[col2 * y + x] = n[col2 * y + x] + l[col1 * y + k] * m[col2 * k + x];
    }
}

int main(){
    // # int row1 = 2, row2 = 3, col1 = 3, col2 = 2;
    int a[row1][col1] = {{1, 2, 3}, {4, 5, 6}};
    int b[row2][col2] = {{9, 8}, {6, 5}, {3, 2}};
    int c[row1][col2];
    int *d, *e, *f;
    int i, j;
    cudaMalloc((void **)&d, row1*col1*sizeof(int));
    cudaMalloc((void **)&e, row2*col2*sizeof(int));
    cudaMalloc((void **)&f, row1*col2*sizeof(int));
    cudaMemcpy(d, a, row1*col1*sizeof(int), cudaMemcpyHostToDevice);
    cudaMemcpy(e, b, row2*col2*sizeof(int), cudaMemcpyHostToDevice);
    dim3 grid(col2, row1);

```

```

    # /* Here we are defining two dimensional Grid(collection of blocks) structure. Syntax is dim3
    grid(no. of columns,no. of rows) */
    matproduct<<<grid,1>>>(d,e,f);
    cudaMemcpy(c, f, row1*col2*sizeof(int), cudaMemcpyDeviceToHost);
    printf("\nProduct of two matrices:\n ");

    for(i = 0; i < row1; i++){
        for(j = 0; j < col2; j++){
            printf("%d\t",c[i][j]);
        }
        printf("\n");
    }
    cudaFree(d);
    cudaFree(e);
    cudaFree(f);
    return 0;
}

```

#### **Matrix Multiplication using CUDA C Output:**

```

Product of two matrices:
 30    24
 84    69

```